

Backpropagation

Felix Portillo

EEC 289Q

Backpropagation (HW3)

- Forward Pass:
 - Inputs: Image data
 - Outputs: $hAct\{L\}$, $L = \text{numHidden}+1$
- Compute cost:
 - $out = \log(hAct\{L\})$
 - Get index of actual label (i.e. index=4 if digit is 3), maybe using 'sub2ind'
 - Compute cost:

$$J(\theta) = - \left[\sum_{i=1}^m \sum_{k=1}^K 1 \{y^{(i)} = k\} \log \frac{\exp(\theta^{(k)\top} h_{W,b}(x^{(i)}))}{\sum_{j=1}^K \exp(\theta^{(j)\top} h_{W,b}(x^{(i)}))} \right] \Rightarrow -\text{sum}(out(index))$$

Backpropagation (cont.)

- From [MultiLayerNeuralNetworks](#):

1. Perform a feedforward pass, computing the activations for layers L_2, L_3 , and so on up to the output layer L_{n_l} .

2. For each output unit i in layer n_l (the output layer), set

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \frac{1}{2} \|y - h_{W,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)})$$

3. For $l = n_l - 1, n_l - 2, n_l - 3, \dots, 2$

For each node i in layer l , set

$$\delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} W_{ji}^{(l)} \delta_j^{(l+1)} \right) f'(z_i^{(l)})$$

4. Compute the desired partial derivatives, which are given as:

$$\frac{\partial}{\partial W_{ij}^{(l)}} J(W, b; x, y) = a_j^{(l)} \delta_i^{(l+1)}$$
$$\frac{\partial}{\partial b_i^{(l)}} J(W, b; x, y) = \delta_i^{(l+1)}.$$

Backpropagation (cont.)

- From MultiLayerNeuralNetworks:

1. Perform a feedforward pass, computing the activations for layers L_2 , L_3 , and so on up to the output layer L_{n_l} .

2. For each output unit i in layer n_l (the output layer), set

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \frac{1}{2} \|u - h_{W,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)})$$

3. For $l = n_l - 1, n_l - 2, \dots, 2$,

For each node i in layer l , set

$$\delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} W_{ji}^{(l)} \delta_j^{(l+1)} \right) f'(z_i^{(l)})$$

4. Compute the desired partial derivatives, which are given as:

$$\frac{\partial}{\partial W_{ij}^{(l)}} J(W, b; x, y) = a_j^{(l)} \delta_i^{(l+1)}$$

$$\frac{\partial}{\partial b_i^{(l)}} J(W, b; x, y) = \delta_i^{(l+1)}.$$

$$\delta^{(n_l)} = - \sum_{i=1}^m \left[\left(1\{y^{(i)} = k\} - P(y^{(i)} = k | x^{(i)}; \theta) \right) \right]$$

What you need

- Get one-hot encoding of actual label
 - Ex. `y_actual = zeros(size(hAct{L}));` `y_actual(index) = 1;`
- For each layer, need to compute ∂ , ∇W , ∇b
- Also need two more variables within algorithm, we'll call them dz and g'
- Store gradients for W , (i.e. ∇W) in `gradStack{l}.W` and b (i.e. ∇b) in `gradStack{l}.b`

The algorithm

- Compute forward pass and store outputs of each layer in $\text{hAct}\{l\}$, where l is the layer number.
- Compute δ $\delta^{(n_l)} = -(y - a^{(n_l)}) \bullet f'(z^{(n_l)})$
 - For softmax output: $\delta = \text{hAct}\{L\} - y_{\text{actual}}$;
- For layer $l = L:-1:1$

if $l = L$, is last layer

$g' = \text{ones}(\text{size}(\delta))$

else

$g' = \text{derivative of activation function (i.e. for sigmoid: } g' = \text{hAct}\{l\} * (1 - \text{hAct}\{l\}) \text{)}$

$\text{dz} = \delta * g'$ %update dz using current δ value, which is error from previous layer

if $l > 1$, is not first layer %Store weight gradients

$\text{gradStack}\{l\}.W = \text{dz} * \text{hAct}\{l-1\}^T$ % use the activation of previous layer

else, is first layer

$\text{gradStack}\{l\}.W = \text{dz} * \text{data}^T$ % use input image

$\text{gradStack}\{l\}.b = \text{sum}(\text{dz})$; %store bias gradients

$\delta = \text{stack}\{l\}.W^T * \text{dz}$; %update δ for next iteration of loop


$$\delta^{(l)} = \left((W^{(l)})^T \delta^{(l+1)} \right) \bullet f'(z^{(l)})$$

$$\begin{aligned} \nabla_{W^{(l)}} J(W, b; x, y) &= \delta^{(l+1)} (a^{(l)})^T, \\ \nabla_{b^{(l)}} J(W, b; x, y) &= \delta^{(l+1)}. \end{aligned}$$

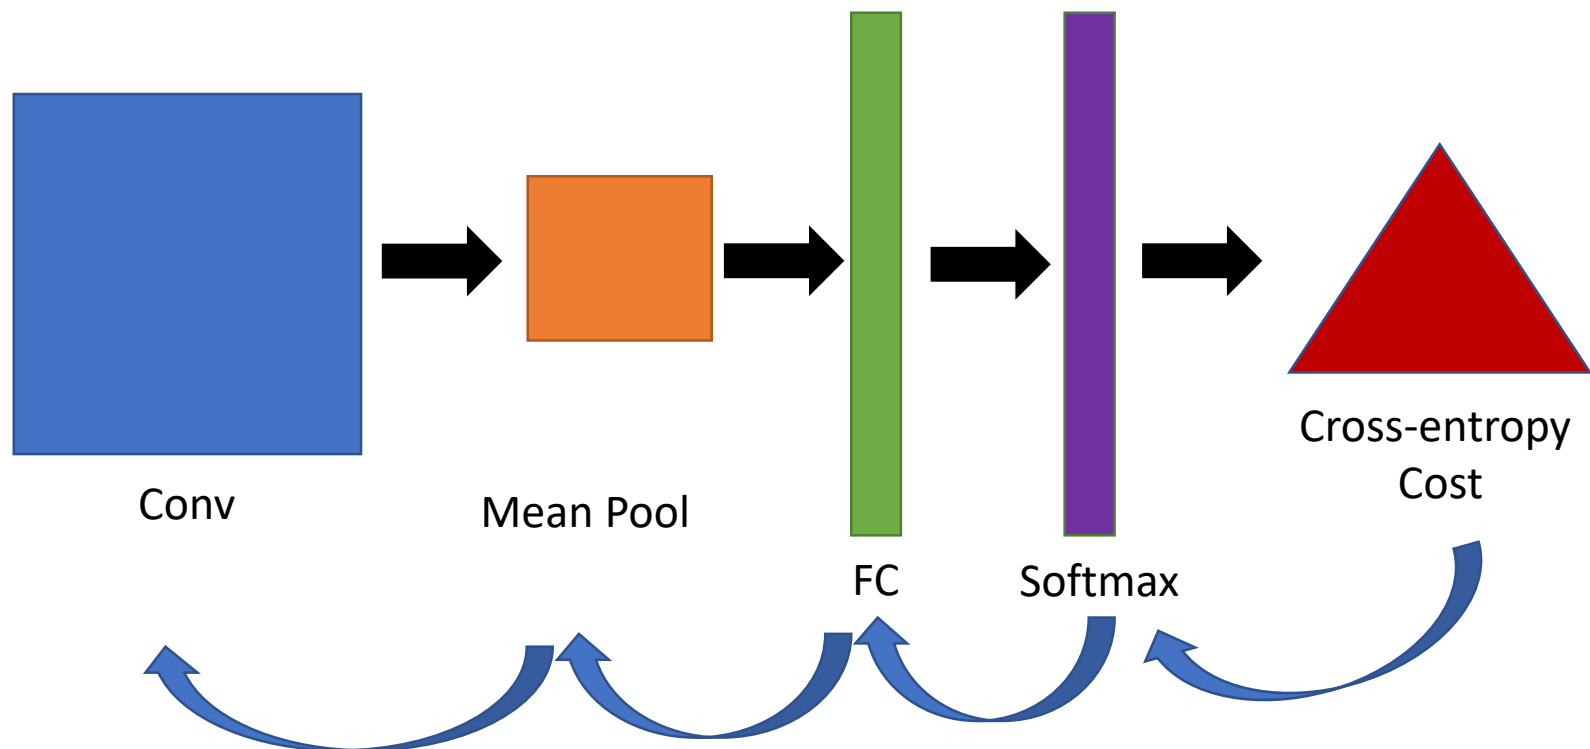
Things to remember

- At end, 'supervised_dnn_cost.m' returns 'grad' so use:
 `[grad] = stack2params(gradStack)`
- May be easier to do vectorized implementation over batch for speed and to avoid confusion.
 - Double check sizes of each variable to ensure they are compatible!


HW4



HW4



HW4



Key Points and Tips

- Backprop has same general concepts for both HW3 and HW4
 - Calculate error, δ , at last layer(L) (softmax error)
 - Go through each layer in network to compute δ for layers starting at L-1
 - Find derivative of activation at each layer, if there is activation at that layer
 - Multiply error from layer l+1 with weights of current layer and/or derivative of activation
It depends on if it's an FC layer or conv layer, you just need to follow directions according to operation. (Details specified on [website](#) for conv)
 - Compute gradients for weights and biases
 - For FC layers, follow outline from HW3, otherwise do conv gradient calculation based on website.
 - Ex. Compute gradients using inputs into layer using either multiplication (FC layer) or convolution with rotated error (conv layer)

Key Points and tips

- Difference b/w backprop for HW3 and HW4
 - HW3: Computes errors (∂) and gradients within same loop
 - HW4: May be best to compute errors (∂) and gradients separately since the layer types are different. Meaning, don't use a loop to go backwards through each layer.
- Remember to reshape matrix after mean pooling into a vector for the FC Layer
 - Similarly, remember to reshape the error vector from the pooling back into a matrix before applying the 'kron' function.
- Pooling has no parameters to update so no need to find gradient for that layer.
- In the end, only returning gradients for W_{conv} , b_{conv} , W_{FC} , b_{FC}

Key Points and Tips (SGD)

- For SGD function, remember to check the order the parameters are initiated in 'cnnInitParams' to help with the theta update.
 - Hint: Makes it easier for vectorized weight updates